

Contact

vasutejas123@gmail.com

www.linkedin.com/in/tejasvasudev
(LinkedIn)

Top Skills

API Development

Amazon Web Services (AWS)

Databricks

Tejas Vasudev

Software Engineer III at JP Morgan

Jersey City, New Jersey, United States

Summary

Software Engineer at JP Morgan. Graduated from Penn State University with a B.S. in Computer Science.

Experience

JPMorgan Chase & Co.

Software Engineer III

September 2023 - Present (3 years 1 month)

Jersey City, NJ

- Led the architecture and delivery of a Databricks-AWS Glue integration that automated governed data access across multiple AWS account sharing models, reducing storage costs and driving 10,000+ user onboardings to the platform.
- Designed a Lambda-based provisioning workflow that validates requests, orchestrates approvals and provisions AWS (IAM roles, Lake formation grants, RAM shares) and Databricks resources across multiple AWS account-sharing models.
- Exposed APIs that allow users to onboard additional databases and tables, manage grants, and assign users/groups to Db/Tables for fine-grained role-based access controls.
- Owned the initiative end-to-end, from requirements gathering and POC through architecture, cross-team alignment, rollout and production adoption.
- Built an approval-driven provisioning workflow that automated Databricks external location and catalog management, turning a multi-step manual onboarding process into self-serve API calls, reducing SRE support tickets by 90%, and scaling onboarding with no added operational overhead
- Redesigned a complex grants workflow by eliminating error-prone state reconciliation logic and establishing Databricks as the system of record, reducing endpoint latency, storage costs and data duplication while cutting grants-related support tickets by 80%.

- Introduced automated integration testing into the CI/CD pipeline and drove its inclusion into our team's SDLC, reducing production incidents and improving release confidence.

Capital One

3 years 2 months

Senior Associate Software Engineer | Monitoring & Observability

August 2021 - September 2023 (2 years 2 months)

San Francisco, CA

- Designed and implemented a scalable backend service for an enterprise reliability and compliance dashboard, optimizing data aggregation and low-latency retrieval to support thousands of daily users and establishing the firm-wide standard for reliability and compliance reporting.
- Built a scalable request-processing layer that filtered redundant deployment events before they reached the core monitoring service, reducing request volume by 70% and enabling enterprise-wide adoption.
- Identified a scalability risk in the onboarding workflow and drove the design of an automated New Relic account-discovery service, reducing operational overhead, eliminating recurring customer onboarding issues, and improving platform readiness for broader adoption.

Software Engineer

August 2020 - August 2021 (1 year 1 month)

- Developed a POC Full Stack app using Angular, Python Flask, Postgres DB, and included functionality to scan business cards (Textextract OCR & ML Model for entity recognition). This helps our Sales team to manage contacts and build relationships with Customers.
- Worked on a ETL Java framework that aggregates data from Salesforce, processes it, and outputs to Data Lake. As part of a migration activity, deployed code changes to register and parse data against a newer internal API, and resolved compliance issues.
- Implemented a monitoring tool for our data pipelines that uses Python Airflow jobs to compare destination data against the source. The email alerts sent out helps our team gain awareness into failures early on.

Capital One

Software Engineering Intern

June 2019 - August 2019 (3 months)

Dallas-Fort Worth Area

Implemented real time reporting platform which enables Auto Finance Managers to generate reports 90% faster than current architecture.

- Reduced end to end data read, compute, and write time from 10 minutes to 52 seconds by refactoring and optimizing slow Golang plugins into efficient SQL scripts, and implementing proof of concept data architecture.
- Debugged issues related to productionizing architecture where data consumer (written in Golang and hosted on Marathon) was unable to periodically write data into MemSql DB.
- Validated platform results by comprehensively testing against Golang plugins to ensure data consistency for end users.

Cisco

Software Engineering Intern

May 2018 - August 2018 (4 months)

San Francisco Bay Area

Developed performance analytics and report generation tool for Intersight (Cisco Cloud based server management solution)

- Programmed python script to retrieve, aggregate, and filter information from Intersight. Then fed this data into Jupyter Notebook data Science tool to perform additional customization and generate concise reports.
- Worked closely with DevOps team to target and eliminate security concerns before generating reports against production data.
- Generated reports gave Engineering and TAC crucial information which aids in debugging, also allows Sales/Marketing to gain a better understanding of real-time product performance.

Education

Penn State University

Bachelor of Science (B.S.), Computer Science · (2016 - 2020)

Cupertino High School

High School Diploma · (2014 - 2016)